



Research Paper

A regime-switching factor model for mean–variance optimization

Giorgio Costa and Roy H. Kwon

Department of Mechanical & Industrial Engineering, University of Toronto, 5 King's College Road, Toronto, ON M5S 3G8, Canada; emails: gcosta@mie.utoronto.ca, rkwon@mie.utoronto.ca

(Received January 19, 2019; revised and accepted August 21, 2019)

ABSTRACT

We formulate a novel Markov regime-switching factor model to describe the cyclical nature of asset returns in modern financial markets. Maintaining a factor model structure allows us to easily derive the asset expected returns and their corresponding covariance matrix. By design, these two parameters are calibrated to better describe the properties of the different market regimes. In turn, these regime-dependent parameters serve as the inputs during mean–variance optimization, thereby constructing portfolios adapted to the current market environment. Through this formulation, the proposed model allows for the construction of large, realistic portfolios at no additional computational cost during optimization. Moreover, the viability of this model can be significantly improved by periodically rebalancing the portfolio, ensuring proper alignment between the estimated parameters and the transient market regimes. An out-of-sample computational experiment over a long investment horizon shows that the proposed regime-dependent portfolios are better aligned with the market environment, yielding a higher ex post rate of return and lower volatility than competing portfolios.

Keywords: asset allocation; Markov regime switching; factor model; mean–variance optimization; robust optimization.

1 INTRODUCTION

The application of mathematical optimization to portfolio construction has been prevalent in both academia and industry since the introduction of modern portfolio theory (MPT) in Markowitz (1952). The mean–variance optimization (MVO) problem introduced by MPT analyzes the trade-off between risk and return to construct so-called optimal portfolios. These portfolios provide an optimal balance of risk and return relative to an investor’s risk appetite. However, while MVO has received widespread attention as a quantitative tool in asset management, it has also been criticized for its susceptibility to estimation error in its input parameters, namely the asset expected returns and covariance matrix. The discrepancies between the estimated parameters and their ex post realizations can severely hinder the performance of an “optimal” portfolio if its optimality is determined solely through these estimates. The sensitivity of MVO to its inputs has been extensively explored in the literature, with important examples shown in Best and Grauer (1991), Merton (1980) and Broadie (1993). A widely referenced conclusion found in Chopra and Ziemba (1993) states that estimation errors during optimization can lead to overconcentrated portfolios, arguing that errors in expected returns can have a ten times larger impact than errors in the covariance matrix.

Measurement of parameter uncertainty and its effect during mathematical optimization is a discipline in its own right and has found several applications in MVO. For example, a heuristic approach known as portfolio resampling (Michaud and Michaud 2008) uses a Monte Carlo method to mitigate the effect of estimation noise during optimization. Robust optimization methods have been used to introduce uncertainty sets around the estimated parameters in order to explicitly capture this estimation error during optimization (Goldfarb and Iyengar 2003; Tütüncü and Koenig 2004), and, more recently, to introduce uncertainty regarding the choice of probability distribution governing the asset returns (Delage and Ye 2010). Portfolio optimization can be enhanced not only by utilizing robust techniques, but also by focussing on the source of uncertainty: the input parameters and the methods by which these are estimated. One example is the application of shrinkage estimators to improve the quality of estimated parameters. Shrinkage can be a particularly useful tool during portfolio optimization when insufficient data is available to produce reliable estimates (Jorion 1986; Ledoit and Wolf 2003).

Following a similar rationale, the aim of this paper is to improve the estimation of input parameters for portfolio optimization. To achieve this, we introduce a novel regime-switching factor model of asset returns that will allow us to incorporate an additional dimension of risk into the estimated parameters. Before delving into this topic, we first discuss the financial application of generic factor models and then regime-switching models.

Factor models attempt to explicitly explain the behavior of a random variable either through a single factor, such as the capital asset pricing model (Lintner 1965; Mossin 1966; Sharpe 1964), or through a combination of multiple factors, such as the Fama–French three-factor model (Fama and French 1993). Factor models have become popular in finance due to the economic relevance of the factors and their ability to explain and quantify different sources of an asset’s systematic risk. One important application of these models is to derive the estimated asset expected returns and covariance matrix, where the inherent properties of the factors are used to explain the asset returns and volatility.

Our intention is to maintain a traditional factor model structure that will allow us to naturally derive these two estimated parameters, while also introducing a regime-switching framework to capture the cyclical nature of financial markets. The original application of regime-switching in finance, proposed in Hamilton (1989), was to describe the abrupt changes observed in business cycles. Due to their similarities, this concept was naturally extended to financial markets to describe the transitions from periods of stability and growth (typically known as bullish market regimes) to periods of financial distress (typically known as bearish market regimes). The literature suggests a two-regime framework can effectively describe this cyclical behavior (Ang and Bekaert 2002a,b; Costa and Kwon 2019; Levy and Kaplanski 2015; Nystrup *et al* 2018). However, there is no consensus on the correct number of regimes. Both Guidolin and Timmermann (2007) and Bae *et al* (2014) identify four separate market regimes to explain the joint distribution of asset returns.

This paper maintains a purely data-driven approach. Since the overall signal-to-noise ratio of the financial market returns is low, our preference is to employ a two-regime framework. The changes in regime can be described through a Markov chain, a property that we exploit in the design of our proposed regime-switching factor model. A Markov chain that describes the transition between latent regimes is sometimes referred to as a hidden Markov model (HMM). Indeed, all the aforementioned literature uses HMMs to model the cyclical behavior of financial markets (a comprehensive literature review on the use of regime-switching models in empirical finance is provided in Guidolin (2011)). In particular, our paper follows the Markov process described in Hamilton (2010) and Ang and Timmermann (2012). The Markov process is used to describe the regime dependency of the factors that drive our factor model of asset returns, allowing us to preserve the structure of a traditional linear model such as the Fama–French three-factor model (Fama and French 1993).

Our proposed factor model has the following rationale. If the factors that drive the asset returns are regime dependent, then the assets themselves must be governed by these market regimes. In Costa and Kwon (2019), the factor model assumes that the factors alone can fully explain the regime dependency of the assets, that is, this behavior can be explained systematically. Instead, in a similar fashion to Liu and

Chen (2016), this paper proposes a factor model structure that allows us to capture both the systematic and idiosyncratic regime dependency of the financial assets.

Through this regime-switching factor model, we can derive estimates of the expected returns and a covariance matrix that both implicitly incorporate a regime dependency, allowing for a more faithful representation of the asset behavior under the current market regime. This model follows a similar framework to that in Costa and Kwon (2019), where the underlying assumption is that the effect of the regime changes on the assets could be fully explained by the factors alone. However, the model in this paper builds on that assumption, and instead assumes that both the idiosyncratic and systematic behavior of the assets is regime dependent. Subsequently, the regime-dependent asset expected returns and covariance matrix can be used to construct MVO portfolios that are better adapted to the current market environment.

The result is a single-period regime-dependent portfolio tailored to the current market regime. “Single-period” denotes the static nature of traditional MVO, where we find an optimal solution today to be used for all subsequent time periods without consideration of any time dependency observed by our input parameters or our decision variable. HMMs often align well with existing dynamic programming models for asset allocation in a multiperiod setting (Ang and Bekaert 2002a; Nystrup *et al* 2018). An alternative is to use scenario-based stochastic programming to incorporate the uncertainty arising from the regime-switching model, as presented in Bae *et al* (2014). However, we favor a single-period approach to improve computational tractability and scalability, avoiding the “curse of dimensionality” often encountered in dynamic programming (Bellman 1961). In other words, a single-period approach allows us to efficiently find optimal portfolios with a large basket of assets. Further, this static model is enhanced by the transient nature of a regime-switching model, reflecting the nonstationarity of the estimated asset parameters. This property can be further exploited through periodic re-estimation of parameters and rebalancing of the portfolio, allowing the portfolio to align itself with the market over a long investment horizon. This, in turn, is compatible with traditional asset management practice, where institutions rely on rebalancing policies with fixed calendar intervals, eg, quarterly, semiannually or annually (Michaud and Michaud 2008).

The experimental results in Section 4 highlight the advantages of a regime-dependent portfolio. Given that estimation errors in asset expected returns make a larger contribution than errors in the covariance matrix (Chopra and Ziemba 1993), it comes as no surprise that any method that improves the quality of estimates can lead to improvements in out-of-sample MVO portfolio performance. That said, our experimental results show a regime-dependent portfolio is able to deliver higher risk-adjusted returns over a long investment horizon, consistently outperforming other portfolios. Another interesting result is given by the minimum-variance portfolio,

demonstrating that, even when expected returns are ignored, the proposed regime-switching framework is able to deliver a better estimate of the covariance matrix, highlighted by having lower ex post volatility. This corroborates the findings of Costa and Kwon (2019), where the use of a regime-switching framework was found to improve the performance of risk-parity portfolios, which are also independent of asset expected returns.

This paper introduces a regime-switching factor model that aims to seamlessly integrate the cyclic property of the market into single-period portfolio optimization while remaining highly tractable and scalable. Moreover, we can overcome the undesired intrinsic static nature of MVO through periodic rebalancing, allowing the portfolio to adapt to current market conditions and further exploiting the benefits of the regime-switching framework. To the best of our knowledge, this regime-switching factor model and its application in MVO are novel.

1.1 Outline of the paper

Section 2 introduces a Markov regime-switching factor model of asset returns. This model allows us to derive the regime-dependent expected return and covariance matrix, implicitly capturing the market dynamics. Section 3 shows how these parameters can be utilized in MVO and its variant, the minimum-variance problem. In addition, we discuss the robust MVO formulation from Goldfarb and Iyengar (2003), which serves as a competing model for comparison. The results of a numerical experiment are presented in Section 4, where we implement and compare the different models under three different rebalancing policies. Section 5 summarizes the findings and contribution of this paper.

2 REGIME-SWITCHING FACTOR MODEL

In this section, we introduce a Markov regime-switching factor model of asset returns. Let us start with a generic factor model. Suppose that the random returns of n assets can be explained through a linear combination of m explanatory factors. Thus, we have that

$$\mathbf{r}_t = \boldsymbol{\mu} + \mathbf{V}^T \mathbf{f}_t + \boldsymbol{\varepsilon}_t, \quad (2.1)$$

where $\mathbf{r}_t \in \mathbb{R}^n$ are the asset random returns at time t , $\boldsymbol{\mu} \in \mathbb{R}^n$ are the asset expected returns, $\mathbf{f}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{F}) \in \mathbb{R}^m$ are the centered factor returns at time t , $\mathbf{V} \in \mathbb{R}^{m \times n}$ is the matrix of factor loadings and $\boldsymbol{\varepsilon}_t \sim \mathcal{N}(\mathbf{0}, \mathbf{D}) \in \mathbb{R}^n$ are the residual returns. $\mathbf{F} \in \mathbb{R}^{m \times m}$ and $\mathbf{D} \in \mathbb{R}^{n \times n}$ denote the factor covariance matrix and the diagonal matrix of residual variance, respectively. Note our choice to use centered factor returns, which explain only the variability of the asset returns. Thus, the constant term in

(2.1) is equal to the asset expected returns. Stemming from this factor model, the asset covariance matrix, $\Sigma \in \mathbb{R}^{n \times n}$, can be derived as

$$\Sigma = V^T F V + D. \tag{2.2}$$

Thus, we have that $r_t \sim \mathcal{N}(\mu, \Sigma)$.

This factor model assumes that the residual returns are independent of one another, ie, $\text{cov}(\varepsilon_t^i, \varepsilon_t^j) = 0$ for $i \neq j$; it also assumes that the residual returns are independent of the factor returns, ie, $\text{cov}(\varepsilon_t^i, f_t^j) = 0$ for all i, j . However, this model does not assume the factors are independent of one another, ie, the factor covariance matrix, F , is not required to be a diagonal matrix. In practice, F must be positive semidefinite, ie, $F \succeq 0$.

In practice, the parameters μ , V and D are estimated by linear regression using historical asset and factor returns, while the parameter F is estimated directly from the factor returns. Throughout this process, one important assumption is made: that recent historical data can accurately estimate the future behavior of the asset returns r_t . Now, as an example, imagine we wish to calibrate our factor model at the onset of the financial crisis. Using recent historical data to do this may provide a false sense of stability, given the bullish market observed prior to the crisis. Thus, the factor model in (2.1) would fail to explain the abrupt change in behavior of the asset returns.

Consider a market-dependent factor, such as the “excess market return” factor from the Fama–French three-factor model (Fama and French 1993).¹ We could describe this change in behavior as a transition into a different market regime; this transition can be described through a Markov process (Ang and Timmermann 2012; Hamilton 2010). Thus, we assume the factor itself is governed by a regime-dependent probability distribution that describes the observed factor returns.

As an example, assume two market regimes exist: bullish and bearish. We can use a discrete-time Markov chain to describe the probability that our state variable s_t will be in regime j at time t given that our previous state was $s_{t-1} = i$, independent of any further preceding states, ie,

$$P\{s_t = j \mid s_{t-1} = i, s_{t-2} = k, \dots\} = P\{s_t = j \mid s_{t-1} = i\}.$$

The state variable s_t is latent and corresponds to the market regime at time t . However, in practice, its probability distribution can be estimated from the observable market data through the Baum–Welch algorithm (Baum *et al* 1970). This algorithm is a special case of expectation maximization, an iterative technique used to find the

¹ We note that all three factors in the Fama–French three-factor model are market dependent, but we choose the “excess market return” as the best example among the three factors to represent the market as a whole.

maximum likelihood. In the case of our two-regime example, the transition matrix governing the probability of switching from one regime to another is

$$\Gamma = \begin{bmatrix} \gamma_{11} & \gamma_{12} \\ \gamma_{21} & \gamma_{22} \end{bmatrix}, \quad (2.3)$$

where γ_{ij} is the probability that regime i at time t will be followed by regime j at $t + 1$. Thus, by definition, each row must sum to 1, ie, $\sum_{j=1}^2 \gamma_{ij} = 1$ for $i = 1, 2$.

This Markov chain can be used to explain the behavior of the factors that drive the asset returns. The model in Costa and Kwon (2019) made the underlying assumption that the factors alone can fully explain the regime dependency of the assets. In other words, the regime dependency of the assets was due to the systematic change of the factors. Therefore, the idiosyncratic asset parameters were assumed to remain the same under both market regimes. This assumption is rather restrictive, particularly when modeling the idiosyncratic risk of the assets between shifting regimes. (A brief description of the regime-switching factor model in Costa and Kwon (2019) can be found in Appendix A online.)

We expand on this idea and assume that the idiosyncratic behavior of the assets is also governed by the different market regimes. The factor model proposed in this paper is structured as the superposition of two individual factor models,

$$r_t = I_{1t}(\mu_1 + V_1^T f_{1t} + \varepsilon_{1t}) + I_{2t}(\mu_2 + V_2^T f_{2t} + \varepsilon_{2t}), \quad (2.4)$$

where each factor model corresponds to a market regime. The indicator variable is $I_{it} = 1$ when the current regime is $s_t = i$ and $I_{it} = 0$ if $s_t \neq i$. As per the factor model proposed in Liu and Chen (2016), all the coefficients from regression are assumed to be regime dependent. Thus, we have that $\mu_1, \mu_2 \in \mathbb{R}^n$ are the regime-dependent expected returns, $V_1, V_2 \in \mathbb{R}^{m \times n}$ are the factor loadings under each regime and $\varepsilon_{1t} \sim \mathcal{N}(\mathbf{0}, \mathbf{D}_1) \in \mathbb{R}^n$ and $\varepsilon_{2t} \sim \mathcal{N}(\mathbf{0}, \mathbf{D}_2) \in \mathbb{R}^n$ are the residuals. Moreover, as in Costa and Kwon (2019), the factors themselves are governed by this Markov chain. Therefore, at some point in time t , the factor returns are described as $f_{1t} \sim \mathcal{N}(\mathbf{0}, \mathbf{F}_1)$ if we are under regime 1, or $f_{2t} \sim \mathcal{N}(\mathbf{0}, \mathbf{F}_2) \in \mathbb{R}^m$ if we are under regime 2.

We will use (2.4) as our regime-switching factor model for the remainder of this paper. The use of centered factors provides the added benefit of simplifying the subsequent derivation of the expectation, variance and covariance of the random returns from (2.4), yielding

$$\begin{aligned} \mu_{s_i} &= \gamma_{i1}\mu_1 + \gamma_{i2}\mu_2, \\ \Sigma_{s_i} &= \gamma_{i1}(V_1^T \mathbf{F}_1 V_1 + \mathbf{D}_1) + \gamma_{i2}(V_2^T \mathbf{F}_2 V_2 + \mathbf{D}_2) \\ &\quad + \gamma_{i1}(1 - \gamma_{i1})\mu_1\mu_1^T + \gamma_{i2}(1 - \gamma_{i2})\mu_2\mu_2^T - \gamma_{i1}\gamma_{i2}(\mu_1\mu_2^T + \mu_2\mu_1^T), \end{aligned} \quad (2.5)$$

where $\mu_{s_i} \in \mathbb{R}^n$ and $\Sigma_{s_i} \in \mathbb{R}^{n \times n}$ are the asset expected returns and covariance matrix corresponding to regime i , respectively. In addition, $\gamma_{ij} = \mathbb{E}[I_j \mid s_t = i]$, ie, γ_{ij} is the probability of switching from regime i to regime j . These probabilities correspond to the two-state transition matrix in (2.3).

The benefit of this model is that the cyclical property of the market is captured implicitly through the asset expected returns and covariance matrix in (2.5). These, in turn, can serve as the input parameters during MVO at no additional computational cost, ie, the formulation of this optimization problem is the same as the nominal one.

The parameters in (2.5) lend themselves to single-period portfolio optimization. However, by design, these parameters are calibrated at time t to match the expectation of the current market regime and do not take into consideration any future transitions from one regime to another. Thus, a limitation of this model is the lack of guarantees that our parameters will be properly aligned during any future time periods.

Nevertheless, this limitation can be overcome through periodic re-estimation of the parameters, allowing them to realign with the market. The frequency of re-estimation has a direct impact on the proper alignment of the estimated parameters. However, given the static property of portfolio optimization, this frequency must be reconciled with a realistic portfolio rebalancing policy. In general, this single-period approach is better suited to traditional asset management techniques, where portfolios are periodically rebalanced and there is a long investment horizon. However, this framework is not meant to serve as a self-sufficient, fully automated investment algorithm. Later, in Section 4, we show that any sort of periodic parameter re-estimation and portfolio rebalancing suffices in order to observe the benefits of the proposed model.

2.1 Selection of market regimes

Here, we discuss the number of regimes present in our model. Determining the appropriate number of market regimes is a difficult task. Moreover, estimation procedures, such as the Baum–Welch algorithm (Baum *et al* 1970), often require the user to define the number of regimes a priori. We do not suggest that a two-regime model is the best for all scenarios, but rather its use reflects our preference for parsimony. In direct reference to the argument presented in Ang and Bekaert (2002b), a two-regime specification is the most a data-driven approach can bear without extreme computational problems in estimation, and it suffices to capture the main empirical differences in asset behavior. Other examples of two-regime models can be found in Ang and Bekaert (2002a), Levy and Kaplanski (2015) and Nystrup *et al* (2018).

If desired, an investor with a predisposition toward a specific number of regimes can modify (2.4) to accommodate additional market regimes and re-derive the regime-dependent parameters in (2.5) accordingly. Increasing the number of regimes

TABLE 1 BIC values for Markov models with different numbers of regimes.

	Two regimes	Three regimes	Four regimes
BIC	−1857.0	−1825.7	−1779.9

has both potential benefits and drawbacks. In theory, increasing the number of regimes may allow us to detect additional nuances within the raw market data, providing a better fitting model. However, this requires the estimation of additional parameters, which is computationally more expensive and also increases the potential for estimation error. The latter is particularly important due to the sensitivity of portfolio optimization to error in estimated parameters (Chopra and Ziemba 1993). With this in mind, we note that there is an existing literature advocating a greater number of regimes. Guidolin and Timmermann (2007) and Bae *et al* (2014) are examples of four-regime models.

We proceed to justify our choice of the number of regimes through an empirical analysis. The analysis will rely on two criteria. The first is the Bayesian information criterion (BIC), introduced in Schwarz (1978). Briefly, the BIC is a measure of log-likelihood penalized by the number of estimated parameters. The BIC is expressed as

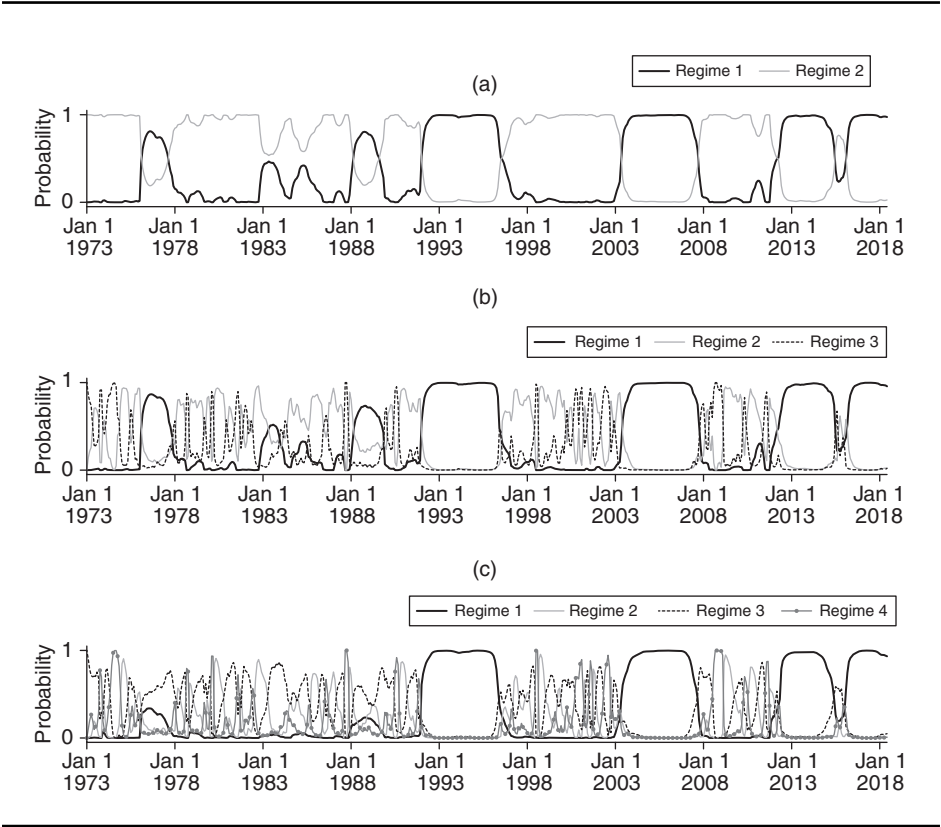
$$\text{BIC} = -2 \ln \hat{L} + (lb + lb^2 + l^2) \ln p, \quad (2.6)$$

where l is the number of regimes, b is the number of factors used for regime detection ($b = 1$ in our case), \hat{L} is the likelihood calculated from the Baum–Welch algorithm and p is the total number of observations in the time series. A lower value of BIC is preferred. The second criterion is more qualitative and relies on the inspection of the plot of smoothed probabilities from the Baum–Welch algorithm.

The analysis is performed on the monthly “excess market return” corresponding to the time period January 1973–June 2018. The data was obtained from Kenneth R. French’s data library (French 2019) and corresponds to the market factor of the Fama–French three-factor model (Fama and French 1993). This is part of the data we will use in Section 4 during the computational experiments. For now, we perform this empirical analysis to justify our choice of regimes.

Table 1 shows that a Markov model with two regimes is the most desirable since it has the lowest BIC. However, we note that the three- and four-regime models are not poor choices either, given that all models have a similar BIC. Thus, we refer to Figure 1 to survey the estimated regimes historically. It is interesting to see that regime 1 is consistent for all three Markov models. However, there is no consistency when we inspect the other regimes. It is the noise observed in the three- and four-regime models that persuades us to use a two-regime model. Again, citing our preference for

FIGURE 1 Estimated regime switches from (a) two-regime, (b) three-regime and (c) four-regime Markov models.



parsimony, a two-regime model suffices to capture the main empirical differences between regimes. Thus, for the remainder of this paper, we assume a two-regime environment.

3 INVESTMENT MODELS

This section presents the two investment models that will be tested under the regime-switching framework: mean–variance optimization and minimum-variance optimization. Starting with the framework presented in Markowitz (1952), the portfolio expected return and variance are

$$\begin{aligned} \mu_p &= \boldsymbol{\mu}^T \mathbf{x}, \\ \sigma_p^2 &= \mathbf{x}^T \boldsymbol{\Sigma} \mathbf{x}, \end{aligned} \tag{3.1}$$

where μ_p is the portfolio expected return, σ_p^2 is the portfolio variance and $\mathbf{x} \in \mathbb{R}^n$ is the vector of asset weights, ie, the proportion of wealth invested in each asset. Here, the portfolio variance is regarded as the financial risk measure. The problem of optimal asset allocation can be solved by finding a solution \mathbf{x}^* that maximizes our utility subject to a set of prescribed constraints. Typical variants of this framework seek to maximize return while limiting our portfolio variance, or to attain an optimal compromise between risk and expected return. However, in this paper we limit ourselves to the version that attempts to minimize portfolio variance subject to a portfolio expected return constraint.

The next two subsections introduce the variants of MVO and minimum-variance optimization that will be implemented during the computational experiment in Section 4. First, we present the basis, or nominal, investment models and briefly introduce a robust variant. Then, we present the regime-dependent investment models. Thus, we have two investment models, and three variants per model: nominal, robust and regime dependent.

3.1 Nominal and robust models

The MVO model seeks to attain an optimal balance between risk and return. In the case of the nominal model, this takes a deterministic approach where the return and risk are perfectly described by the estimated parameters in (3.1).

If we seek to minimize risk while subject to a target return, the MVO model can be written as

$$\left. \begin{array}{ll} \min_{\mathbf{x}} & \mathbf{x}^T \boldsymbol{\Sigma} \mathbf{x} \\ \text{such that} & \boldsymbol{\mu}^T \mathbf{x} \geq R, \\ & \mathbf{1}^T \mathbf{x} = 1, \end{array} \right\} \quad (3.2)$$

where the first constraint ensures we find a portfolio with an expected return that meets or exceeds a predetermined target, $R \in \mathbb{R}$, and the second constraint is the “budget constraint”, ensuring that the entirety of the available portfolio wealth is spent. We use the notation $\mathbf{1}$ to denote a column vector of appropriate size where all entries are equal to 1. We will refer to this model as the nominal MVO problem for the remainder of the paper.

The MVO problem in (3.2) is a typical example of a quadratic optimization problem with linear constraints. Given that the asset covariance matrix is positive semidefinite,² this is a convex quadratic problem. This, in turn, implies the problem can be efficiently solved computationally and that any solution we find within the feasible set is a global optimal solution.

² A useful property of covariance matrixes in finance is their tendency to be positive definite or positive semidefinite.

Selecting an appropriate target return R should be considered carefully. An excessively high target may lead to unreasonable portfolios with extreme long and short positions, or may even lead to infeasibility of the optimization problem itself. A reasonable approach to size the target return is to define it in terms of the asset expected returns $\boldsymbol{\mu}$. For example, we may require a premium $\kappa \in \mathbb{R}$ above the average asset expected return, ie,

$$R = \left(\frac{1 + \kappa}{n} \right) \mathbf{1}^T \boldsymbol{\mu}. \tag{3.3}$$

The choice of an adequate value of κ retains the feasibility of the optimization problem while maintaining a realistic target return.

Next, we introduce the minimum-variance model. This optimization model seeks to minimize the portfolio risk without imposing any conditions on the portfolio expected return. Thus, the only input required is the asset covariance matrix. The nominal minimum-variance problem is

$$\left. \begin{array}{ll} \min_{\mathbf{x}} & \mathbf{x}^T \boldsymbol{\Sigma} \mathbf{x} \\ \text{such that} & \mathbf{1}^T \mathbf{x} = 1. \end{array} \right\} \tag{3.4}$$

In essence, minimum-variance optimization is the same as MVO, except we relax the optimization model by removing the target return constraint, allowing the model to find the optimal portfolio with the lowest attainable variance.

Finally, we consider a robust variant of the MVO and minimum-variance models. The robust optimization framework implemented in this paper is presented in detail in Goldfarb and Iyengar (2003); for brevity, it is not described further. We chose this particular robust framework because it derives the uncertainty structures directly from a factor model. Specifically, these uncertainty structures are built around the estimated returns $\boldsymbol{\mu}$, factor loadings \mathbf{V} and residual variance \mathbf{D} . As with the nominal models, the only difference between the robust MVO and robust minimum-variance model is the inclusion of a target return constraint, subject to the appropriate uncertainty structure. The formulation of the robust optimization models is presented in detail in Appendix B online. These robust variants will serve as an additional benchmark against the proposed regime-dependent models.

3.2 Regime-dependent models

Here, we introduce the regime-dependent investment models. As discussed in Section 2, the cyclical nature of the market is fully captured through the estimated regime-dependent parameters $\boldsymbol{\mu}_{s_i}$ and $\boldsymbol{\Sigma}_{s_i}$ from (2.5). Thus, the regime-dependent MVO variant can easily be attained by replacing the nominal input parameters with

their regime-dependent counterparts, ie,

$$\left. \begin{array}{ll} \min_x & \mathbf{x}^T \boldsymbol{\Sigma}_{s_i} \mathbf{x} \\ \text{such that} & \boldsymbol{\mu}_{s_i}^T \mathbf{x} \geq R_{s_i}, \\ & \mathbf{1}^T \mathbf{x} = 1, \end{array} \right\} \quad (3.5)$$

where the target return $R_{s_i} \in \mathbb{R}$ is based on the regime-dependent estimated returns $\boldsymbol{\mu}_{s_i}$ in a similar fashion to (3.3). We will refer to this model as the regime-dependent MVO model.

We will take a moment to highlight the importance of the target return in (3.5). Similarly to the nominal model, this constraint ensures the portfolio expected return satisfies a predetermined value. However, calculating the target return from the regime-dependent estimated returns $\boldsymbol{\mu}_{s_i}$ sets a target return adapted to the current market environment. In other words, depending on the current market regime, the regime-dependent target return R_{s_i} will be higher or lower than its nominal counterpart, allowing a more realistic expectation of an attainable return on the portfolio.

The final variant we discuss is the regime-dependent minimum-variance model, which, as before, can be attained by removing the target return constraint. Thus, we have

$$\left. \begin{array}{ll} \min_x & \mathbf{x}^T \boldsymbol{\Sigma}_{s_i} \mathbf{x} \\ \text{such that} & \mathbf{1}^T \mathbf{x} = 1. \end{array} \right\} \quad (3.6)$$

We consider the minimum-variance problem in this paper in order to determine whether a regime-switching model can be beneficial even when estimated expected returns are not used as an input during portfolio optimization. The conclusion in Costa and Kwon (2019) indicates it is beneficial in risk-parity portfolios, where expected returns are also not required. Thus, we wish to test whether this conclusion extends to other quantitative investment models.

Mathematically, the regime-dependent optimization models have the same structure as their nominal counterparts. This means we retain the same level of complexity as the nominal models: these are quadratic and convex, guaranteeing a global optimal solution. From a mathematical optimization point of view, both the nominal and regime-dependent models are generic quadratic problems. However, if we look beyond the optimization framework, we can see that the models are fundamentally different, given their input parameters. The regime-dependent models seek to minimize risk for a given market regime, thereby tailoring the resulting portfolio accordingly. This is in direct contrast to the nominal models, which yield a portfolio befitting a noncyclical market.

TABLE 2 List of assets.

Food products	Tobacco	Beer and liquor	Recreation	Printing
Household products	Apparel	Health care	Chemicals	Textiles
Fabricated products	Construction	Steel works	Electrical equipment	Automobiles
Aircraft, ships, rail equipment	Mining	Coal	Oil and gas	Utilities
Communication	Services	Business equipment	Paper	Transportation
Restaurants and hotels	Wholesale	Retail	Financials	Other

We note that both investment models (MVO and minimum-variance optimization) and their three corresponding variants (nominal, robust and regime dependent) allow for the construction of portfolios with short positions, ie, we do not enforce a “long-only” constraint. The regime-switching model in this paper is only concerned with the estimated input parameters. Thus, any typical constraints on the wealth allocation variable \mathbf{x} , such as limiting our exposure to certain industries or enforcing “long-only” positions, are allowed.

4 COMPUTATIONAL EXPERIMENT

In this section, we perform four separate computational experiments to test the out-of-sample performance of the regime-dependent investment models relative to their nominal and robust counterparts. We choose the Fama–French three-factor model (Fama and French 1993) as the basis to test the regime-switching framework due to its popularity as a trustworthy multiple factor model. However, we note that this framework is applicable to any model with market-dependent factors. The tractability and scalability of the proposed framework is tested by using a basket of thirty diverse assets consisting of thirty different industries, which will serve as the constituents of our optimal portfolios. Table 2 shows the thirty industries under consideration. The historical data was obtained from Kenneth R. French’s data library (French 2019).


An overview of the setup of the first two experiments follows. We perform these experiments using historical monthly returns for both assets and factors. We test the performance of both MVO and minimum-variance models under each of the three variants described in Section 3: nominal, robust and regime dependent. Thus, we construct and test six different portfolios. For the nominal and robust models, the asset expected returns and covariance matrix are estimated through an ordinary least

squares regression stemming from the factor model in (2.2), using two years of historical data (twenty-four observations) immediately preceding the investment period. All estimated parameters are recalibrated every time the portfolios are rebalanced. In addition, the robust portfolios are constructed as outlined in Appendix B online using a confidence level of 0.95. The construction of the regime-dependent portfolios is explained below, and the rebalancing policy is explained thereafter.

Estimation of the regime-dependent parameters requires additional historical data to properly calibrate the Markov model. For this experiment, we use historical factor returns starting from January 1973, thirty years before the start of the out-of-sample test period. We apply the Baum–Welch algorithm (Baum *et al* 1970) to find both the transition probabilities and the smoothed probabilities, with a practical implementation of this algorithm derived from Kritzman *et al* (2012). We use the “excess market return” factor as the sole input to the Baum–Welch algorithm. Assuming a two-regime model, the smoothed probabilities serve to differentiate between bullish and bearish historical periods. This data-driven approach allows us to select periods of historical factor and asset returns corresponding to each regime. For consistency, the most recent twenty-four observations corresponding to each regime are used to perform the regression. We note that these forty-eight observations stretch beyond the immediately preceding two years of historical data and go as far back in time as necessary to ensure the availability of sufficient data for both regimes. We favor this approach to guarantee a consistent number of observations per regime during estimation, improving the stability of the regression model. The appropriate transition probabilities γ_{i1} and γ_{i2} are selected by identifying the “current” regime, ie, by inspecting the most recent observations of the smoothed probabilities and selecting the larger of the two. After the two regime-dependent regression models are calibrated, together with the appropriate transition probabilities, the asset expected returns and covariance matrix are estimated as per (2.5). This process is repeated every time the portfolios are rebalanced. Finally, we note that the regime estimation window, which was originally thirty years long, expands every time we rebalance the portfolio, ie, the Baum–Welch algorithm is applied over an increasingly long time period as the experiment moves forward through time.

The out-of-sample experiments are performed over the time period from January 2003 to June 2018. We conduct three individual experiments over this period, each with a different rebalancing policy, where the portfolios are rebalanced every three, six or twelve months. Our choice of rebalancing frequency is motivated by the practical choices made by investors (Michaud and Michaud 2008). For all three experiments, the in-sample calibration window is rolled forward every time a portfolio is rebalanced in order to use the most recently available data. The experimental results are given in the following format. First, we present an example of the estimated regime changes. Next, we present the results of the MVO portfolios

FIGURE 2 Estimated regime switches over the calibration window (a) before the first investment period and (b) before the last investment period.



Regime 1 corresponds to a bullish market.

under the three different rebalancing policies. Finally, we present the results of the minimum-variance portfolios.

All computations were performed on an Apple MacBook Pro computer (2.8 GHz Intel Core i7, 16 GB 2133 MHz DDR3 RAM) running OS X “High Sierra”. All computations were performed using the Julia programming language (version 1.1.0), using the optimization modeling language JuMP (Dunning *et al* 2017) with IPOPT (version 3.12.6) as the optimization solver.

An example of the estimated changes in regime from this experiment is presented in Figure 2, which plots the smoothed probabilities of a two-regime model prior to the start of the first investment period and prior to the start of the last investment period. The assumption that the factors within the factor model are subject to the market’s cyclical behavior is only valid if at least one of the factors is market dependent. In our experiment, all three Fama–French factors are market dependent. However, we assume the “excess market return” factor conveys the latent market information more faithfully. Thus, the Baum–Welch algorithm is applied only on the observations corresponding to this factor. Using a single factor during the estimation process is less computationally expensive and tends to reduce the level of noise during the estimation process. Nevertheless, this data-driven process may sometimes result in unusual

behavior. For example, a bullish market regime is observed to dominate throughout the first estimate in Figure 2(a). However, this same period (1973–2002) appears to be considered mostly bearish when we include data from 1973 to 2018 (Figure 2(b)). Since, by design, we have imposed a two-regime model, the observed probability distribution that governs a specific regime may change as new data becomes available, ie, a period that was once considered to be bullish may be considered bearish if new data exacerbates the differences between the two regimes.

4.1 MVO experiments

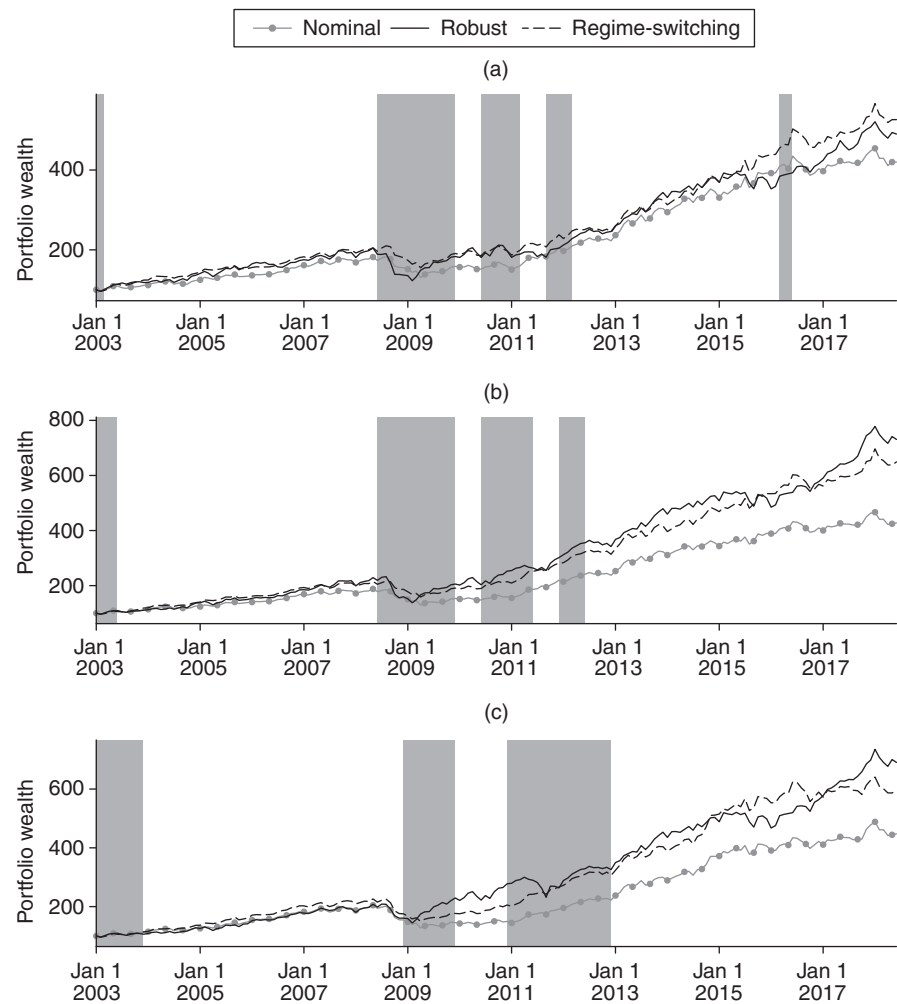
The nominal, robust and regime-dependent MVO portfolios were tested under three different rebalancing policies to observe whether the directional insight of the regime-switching framework can provide any significant benefit over both the nominal and robust portfolios.

The wealth evolution of the MVO portfolios is shown in Figure 3. Within the figure, a shaded or white background indicates whether the regime-switching portfolio was calibrated to align with a bearish or bullish market regime, respectively. We note that both the white and shaded periods are chosen on an *ex ante* basis by the Baum–Welch algorithm, ie, the model estimates the “current” regime using only data from the calibration period during each rebalancing point and is blind to any future data. The resulting estimated regimes align well with observed historical periods of market distress, such as the recession at the end of the early 2000s and the financial crisis of 2007–9.

The plots of the three-, six- and twelve-month rebalancing policy in Figure 3 show that both the robust and regime-dependent portfolios are able to attain a higher terminal wealth than the nominal portfolio. However, upon closer inspection, only the regime-dependent portfolio is able to consistently outperform the nominal. In particular, the robust portfolio suffers a sharp downturn around the end of 2008, with the regime-dependent portfolio being able to remain above the nominal during this period. In addition, the robust portfolio exhibits a more volatile wealth path, oscillating above and below the path of the regime-dependent portfolio.

As discussed in Section 3, another attribute of the regime-dependent MVO model is the selection of a target return. The target return imposes a constraint before optimization takes place, with a lower target relaxing the optimization problem. For this experiment, we ask for a 10% premium over the average asset expected return, ie, $\kappa = 0.1$ in (3.3). Figure 4 shows the difference in the target returns used during optimization of the regime-dependent and non-regime-dependent MVO portfolios, $R_{s_i} - R$, at each rebalancing point. The difference in target return allows us to observe the difference in predisposition of the portfolios toward the market each

FIGURE 3 Wealth evolution of MVO portfolios.

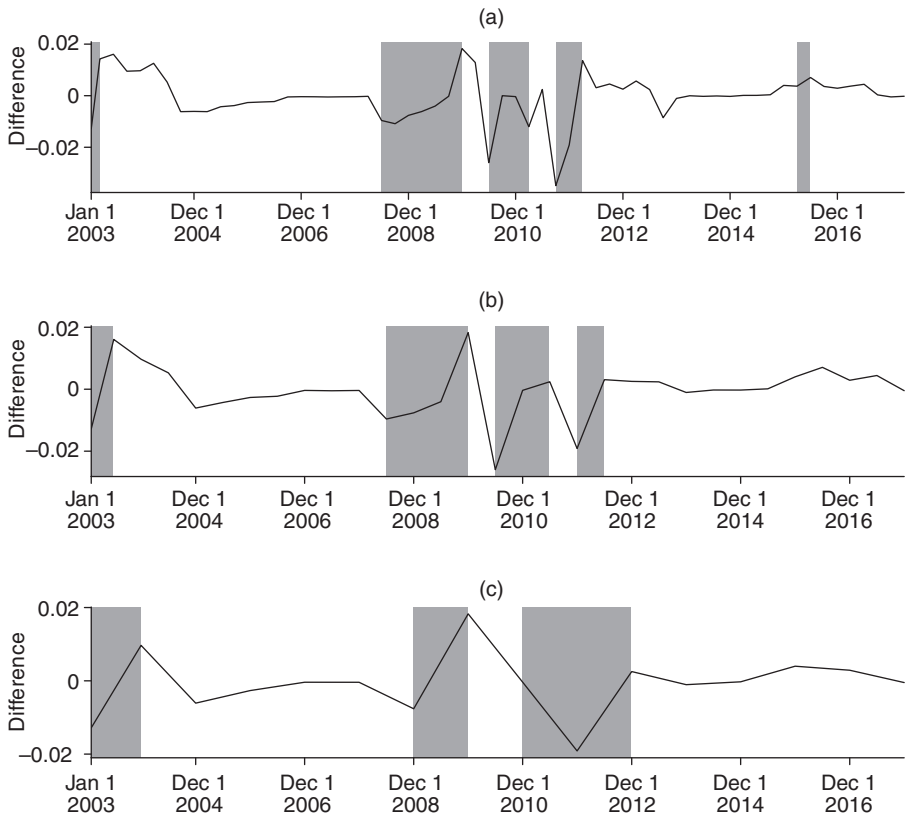


A shaded background indicates time periods where the regime-dependent portfolios were calibrated for a bearish market regime. A white background indicates calibration for a bullish market regime. The wealth evolutions depicted are associated with (a) three-, (b) six- and (c) twelve-month rebalancing policies.

time the portfolios were optimized. From the three plots, it is possible to see that the regime-dependent portfolio has a lower target as it approaches the shaded regions, while it has a higher target as it approaches white regions.

Table 3 summarizes the portfolio performance throughout the entire investment horizon (2003–18) for all three rebalancing policies. These results are calculated

FIGURE 4 Difference in target return between regime-dependent and non-regime-dependent target returns.



Rebalancing period: (a) three months, (b) six months, (c) twelve months.

TABLE 3 MVO portfolios: summary of results.

	Three-month policy			Six-month policy			Twelve-month policy		
	Nom.	Rob.	RS	Nom.	Rob.	RS	Nom.	Rob.	RS
Annual return	0.098	0.109	0.114	0.099	0.138	0.129	0.102	0.134	0.123
Annual volatility	0.111	0.141	0.108	0.110	0.140	0.107	0.116	0.140	0.113
Sharpe ratio	0.881	0.772	1.052	0.900	0.983	1.211	0.878	0.953	1.083
Average turnover	1.145	1.983	1.325	1.632	2.702	1.717	2.261	3.343	2.115

Nom, nominal. Rob, robust. RS, regime switching.

on an ex post basis. From the table, we can see that the regime-dependent MVO portfolio has a consistently better performance than the nominal portfolio in terms of both risk and reward. On the other hand, the robust MVO portfolio exhibits mixed results. The robust portfolio has a high rate of return but maintains an equally high volatility, eroding its risk-adjusted returns. The risk-adjusted returns are given by the corresponding ex post Sharpe ratios (Sharpe 1994). With this in mind, we can see that the regime-dependent portfolio was able to deliver higher risk-adjusted returns over the entire investment horizon, consistently outperforming both the nominal and robust models. The average turnover, which we use to represent potential transaction costs, shows the regime-dependent portfolio is able to maintain a comparable turnover to its nominal counterpart throughout all three rebalancing policies.

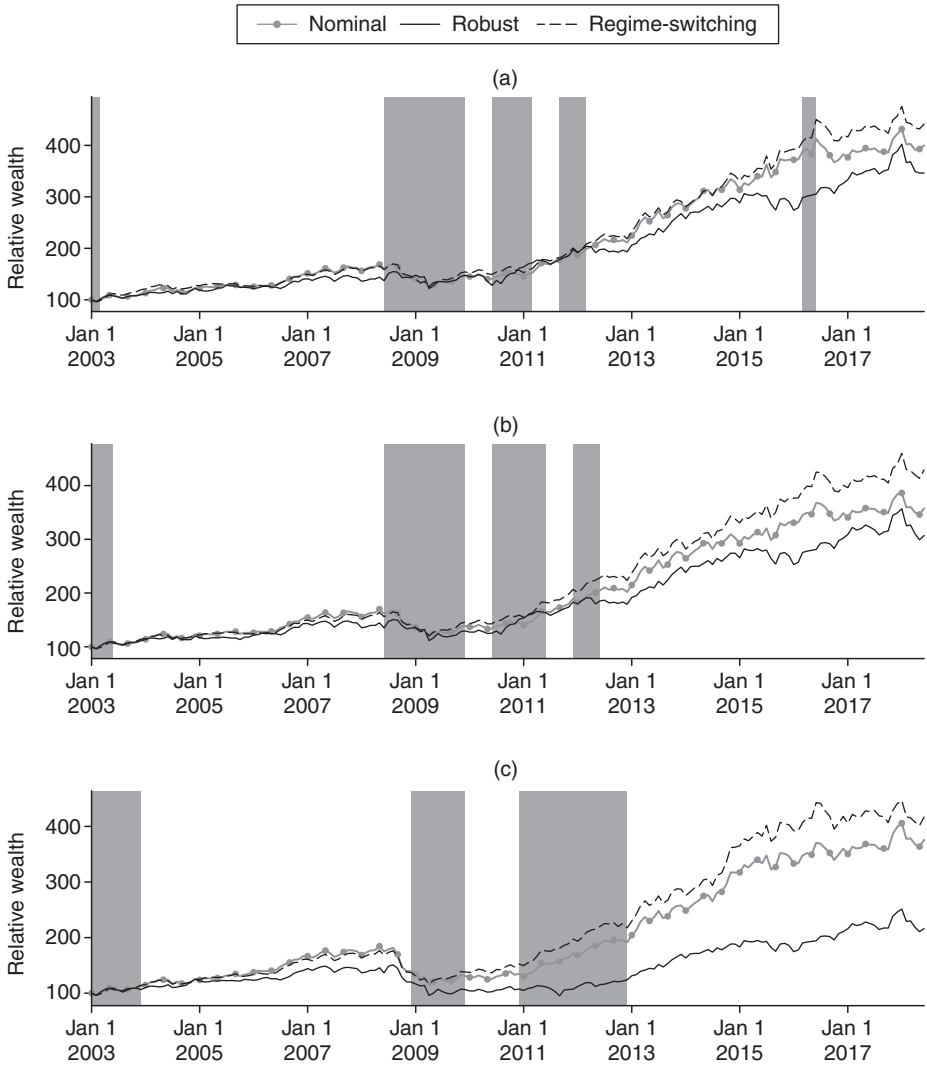
4.2 Minimum-variance experiments

The inclusion of the minimum-variance portfolios in our computational experiments influences whether the regime-switching framework has a measurable effect on investment models that do not require estimated expected returns as an input parameter. In this sense, these computational experiments could corroborate the findings in Costa and Kwon (2019), where a similar regime-switching framework is used in risk-parity portfolios, which are independent of estimated expected returns.

Figure 5 shows the wealth evolution of the minimum-variance portfolios. As before, the shaded and white backgrounds indicate the predisposition of the regime-dependent portfolio to a bearish or bullish market, respectively. For all three rebalancing policies, the wealth path of the regime-dependent model steadily outpaces both its nominal and robust counterparts. This is particularly interesting because these results are fully driven by the estimated covariance matrixes. In particular, the nominal and regime-dependent portfolios closely mimic each other during the first half of the investment horizon, and the regime-dependent model is only able to gain a measurable lead after the financial crisis of 2008. This effect is more pronounced as the rebalancing window increases. Moreover, the three plots suggest the robust minimum-variance model has an excessively conservative outlook, yielding meager results.

Finally, Table 4 summarizes the portfolio performance for all three rebalancing policies. By design, the rate of return attained by minimum-variance portfolios is not driven by any target preset during optimization. Thus, it is difficult to attribute a higher rate of return directly to having better quality estimates of the covariance matrix. Nevertheless, this higher rate of return of the regime-dependent portfolio may be indirectly driven by the lower volatility observed in all three experiments, sheltering the portfolio from undesired downward movements. The ex post Sharpe

FIGURE 5 Wealth evolution of minimum-variance portfolios.



A shaded background indicates time periods where the regime-dependent portfolios were calibrated for a bearish market regime. A white background indicates calibration for a bullish market regime. Rebalancing period: (a) three months, (b) six months, (c) twelve months.

ratios yield the same results as those for the MVO portfolios: the regime-dependent portfolio exhibits higher risk-adjusted returns under all three rebalancing policies while maintaining a similar average turnover to its nominal counterpart.

TABLE 4 Minimum variance: summary of results.

	Three-month policy			Six-month policy			Twelve-month policy		
	Nom.	Rob.	RS	Nom.	Rob.	RS	Nom.	Rob.	RS
Annual return	0.094	0.084	0.101	0.086	0.075	0.099	0.090	0.051	0.097
Annual volatility	0.109	0.120	0.105	0.109	0.119	0.104	0.114	0.124	0.110
Sharpe ratio	0.863	0.698	0.963	0.791	0.633	0.950	0.788	0.415	0.881
Average turnover	0.960	1.212	1.037	1.392	1.752	1.410	1.984	2.466	1.835

Nom, nominal. Rob, robust. RS, regime switching.

4.3 Experiments with a turnover limit

These experiments follow the same methodology as our previous experiments, except we now impose a turnover constraint during optimization every time we rebalance our portfolios. As before, we use monthly data, but we only consider a six-month rebalancing policy. A turnover constraint limits how much an investor is able to change their position on their investments, given they already hold a portfolio. The portfolio turnover serves as a proxy of transaction costs, where an investor must pay a fee to buy and sell their assets on an exchange. This fee is usually small and depends on the relative volume as well as the absolute wealth of the portfolio. Thus, it is common for an investor to place a limit on these fees. To avoid making explicit assumptions about the monetary costs of this transaction, we place a limit on the portfolio turnover instead. For brevity, we conduct this experiment on the nominal and regime-switching portfolios and exclude the robust portfolios.


To model the limit on turnover, we add the following constraint to the MVO models in (3.2) and (3.5) as well as the minimum-variance models in (3.4) and (3.5),

$$\|x - x^0\|_1 \leq C,$$

where $x^0 \in \mathbb{R}^n$ is the vector of weights corresponding to our current holdings and has a constant value, $C \in \mathbb{R}_+$ is our limit on the portfolio turnover, and $\|\cdot\|_1$ is the ℓ_1 -norm operator. We note that the nonlinear ℓ_1 -norm operator can be reformulated as a set of linear constraints. These experiments impose three different limits on turnover: $C = 1.25$, $C = 1$ and $C = 0.75$.

Figure 6 shows the relative wealth of the regime-switching portfolios compared with their nominal counterparts for each instance of the turnover limit. The results are displayed as a relative measure since our aim is to evaluate whether the regime-switching portfolios can maintain an advantage over their nominal counterparts, and whether this advantage is consistent over different turnover limits. Both the MVO and

FIGURE 6 Relative wealth evolution of (a) MVO and (b) minimum-variance portfolios with turnover limits.



Plots show the relative wealth of regime-switching portfolios against their nominal counterparts. A shaded background indicates time periods where the regime-dependent portfolios were calibrated for a bearish market regime. A white background indicates calibration for a bullish market regime.

minimum-variance plots show that this is indeed the case, with the regime-switching portfolios having a consistent advantage over most of the investment horizon.

The results summarized in Tables 5 and 6 corroborate our observations. If we look at the performance metrics, we can see the effect of limiting our flexibility to rebalance our portfolios. As the value of C increases, both the annualized average return and Sharpe ratio decrease. Nevertheless, the regime-switching portfolios still maintain their advantage over their nominal counterparts under all scenarios.

4.4 Experiments with weekly rebalancing

These last experiments involve a different data set. As before, we use the Fama–French three-factor model to estimate both the nominal and regime-dependent input parameters, namely the expected returns and covariance matrix. The assets are also

TABLE 5 Summary of results for MVO portfolios with a turnover limit.

	No limit		C = 1.25		C = 1		C = 0.75	
	Nom.	RS	Nom.	RS	Nom.	RS	Nom.	RS
Annual return	0.099	0.129	0.091	0.117	0.088	0.115	0.085	0.093
Annual volatility	0.110	0.107	0.112	0.106	0.112	0.108	0.111	0.113
Sharpe ratio	0.900	1.211	0.818	1.097	0.780	1.064	0.764	0.822
Average turnover	1.632	1.717	1.200	1.237	0.988	1.000	0.750	0.750

Nom, nominal. RS, regime switching.

TABLE 6 Summary of results for minimum-variance portfolios with a turnover limit.

	No limit		C = 1.25		C = 1		C = 0.75	
	Nom.	RS	Nom.	RS	Nom.	RS	Nom.	RS
Annual return	0.086	0.099	0.089	0.101	0.089	0.101	0.086	0.097
Annual volatility	0.109	0.104	0.111	0.105	0.110	0.106	0.109	0.106
Sharpe ratio	0.791	0.950	0.802	0.962	0.809	0.959	0.783	0.913
Average turnover	1.392	1.410	1.145	1.139	0.973	0.964	0.750	0.739

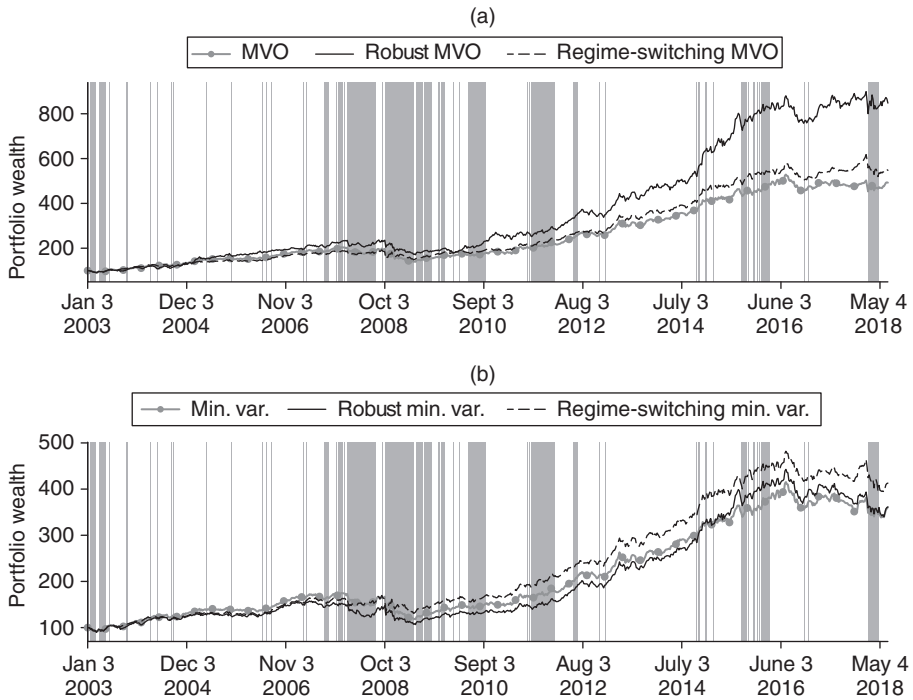
Nom, nominal. RS, regime switching.

the same as before (see Table 2). The time period is also the same, ranging from January 2003 to June 2018. The differences are the frequency of the data and the rebalancing frequency. Instead of monthly returns, we use weekly returns, and we rebalance the portfolios every week (ie, we recalibrate our models and rebalance our portfolios each time a new observation becomes available).

These experiments test the out-of-sample performance of our model in a higher-frequency environment. We note that rebalancing on a weekly basis is an extreme scenario for our proposed model, which is better suited for traditional, long-term portfolio management with periodic rebalancing. Moreover, the lower signal-to-noise ratio of weekly observations also has a negative effect on the data-driven regime calibration through the Baum–Welch algorithm.

The wealth evolution of the MVO and minimum-variance portfolios is shown in Figure 7. From a wealth perspective, the robust model dominates the MVO portfolios. Nevertheless, the regime-dependent MVO portfolio still outperforms its nominal counterpart. Looking at the minimum-variance portfolios, the regime-dependent model maintains a consistent lead for most of the investment horizon. For both MVO and minimum-variance portfolios, these results are quite similar to the previous experiments. The main difference can be seen when we inspect the shaded periods in

FIGURE 7 Wealth evolution of (a) MVO and (b) minimum-variance portfolios with weekly rebalancing.



A shaded background indicates time periods where the regime-dependent portfolios were calibrated for a bearish market regime. A white background indicates calibration for a bullish market regime.

this figure. A weekly rebalancing policy means our regime-dependent portfolios are very sensitive to noise in the estimation of regimes via the Baum–Welch algorithm. This is reflected in the multiple thin, shaded lines.

A summary of results is given in Table 7. The summary statistics show the same overall trend as in the previous experiments, except for the Sharpe ratio of the MVO portfolios. Here, we can see that the robust portfolio returns significantly outperform the competing models, even after adjusting for the volatility incurred. As noted earlier, this suggests that the regime-dependent model suffers from the noise observed in the weekly returns and this is compounded by a weekly rebalancing policy. Nevertheless, the regime-dependent MVO portfolio still retains an advantage over the nominal model. The results of the minimum-variance portfolios align with those of previous experiments, with the regime-dependent model having the best performance of the three competing models.

TABLE 7 Summary of results for portfolios with weekly rebalancing.

	MVO			Minimum variance		
	Nom.	Rob.	RS	Nom.	Rob.	RS
Annual return	0.108	0.137	0.116	0.086	0.086	0.096
Annual volatility	0.125	0.140	0.123	0.123	0.130	0.121
Sharpe ratio	0.864	0.982	0.936	0.699	0.661	0.793
Average turnover	0.238	0.600	0.446	0.158	0.375	0.338

Nom, nominal. Rob, robust. RS, regime switching.

5 CONCLUSION

This paper introduced a regime-switching factor model that reflects the cyclical nature of asset returns in modern financial markets. This model retains the benefits of a traditional factor model, namely the economic relevance of the factors and the ability to easily derive the asset expected returns and covariance matrix. Deriving these parameters through the regime-switching factor model implicitly captures the directional information from the market through a probability-weighted approach, adapting the parameters to the current market regime. In turn, using these parameters during MVO poises the resulting optimal portfolio to take advantage of this directional information, resulting in portfolios with reduced ex post volatility and increased returns. The novelty of the proposed framework is the seamless integration of the regime dependency of the asset returns with the static nature of MVO, reconciling it with the dynamism of the market through periodic portfolio rebalancing.

The regime-switching factor model can easily be derived from any generic model where at least one of the factors observes the cyclical behavior of the market. Explaining the transition between market regimes through a discrete-time Markov chain allows us to retain a data-driven modeling structure where we can use historical observations to calibrate our model and estimate the regime-dependent asset expected returns and covariance matrix. This aligns well with a single-period portfolio optimization model such as MVO. Since regime dependency is already captured through the estimated parameters, we highlight that this framework comes at no additional computational cost during optimization. In general, this single-period approach is well suited for traditional asset management techniques, where portfolios are periodically rebalanced and there is a long investment horizon.

The experimental results show that a regime-dependent MVO portfolio is able to steadily outperform its nominal counterpart over a long investment horizon, exhibiting lower volatility and higher returns. These results are consistent over three different rebalancing policies of varying lengths, showing that regime-dependent

portfolios are quite resilient. Provided some instance of periodic rebalancing takes place, the portfolio is able to adapt to the market regime. Even if there is some regime misalignment between the portfolio and the market through time, the results show that this is outweighed by the benefit of having some partial alignment throughout the investment horizon. Moreover, the proposed model is even shown to be more beneficial than a robust MVO portfolio. The regime-dependent portfolio is able to deliver higher risk-adjusted returns than both the nominal and robust portfolios. These results are sustained even when we do not apply a portfolio target return.

The application of this regime-switching factor model to portfolio optimization opens the door to future lines of research. Extrapolating from the experimental results, a possible improvement is to formulate a robust counterpart of the regime-dependent portfolios by introducing uncertainty structures as in Goldfarb and Iyengar (2003). Motivated by the experimental results in Costa and Kwon (2019), another exciting avenue of future research is the study of rebalancing control policies for single-period portfolio optimization models. This, in turn, could lend itself to the development of a data-driven framework that identifies pertinent market signals, such as changes in regime, to trigger portfolio rebalancing.

DECLARATION OF INTEREST

The authors report no conflicts of interest. The authors alone are responsible for the content and writing of the paper.

ACKNOWLEDGEMENTS

This work was supported by a Mitacs Accelerate grant.

REFERENCES

- Ang, A., and Bekaert, G. (2002a). International asset allocation with regime shifts. *Review of Financial Studies* **15**(4), 1137–1187 (<https://doi.org/10.1093/rfs/15.4.1137>).
- Ang, A., and Bekaert, G. (2002b). Regime switches in interest rates. *Journal of Business & Economic Statistics* **20**(2), 163–182 (<https://doi.org/10.1198/073500102317351930>).
- Ang, A., and Timmermann, A. (2012). Regime changes and financial markets. *Annual Review of Financial Economics* **4**(1), 313–337 (<https://doi.org/10.1146/annurev-financial-110311-101808>).
- Bae, G. I., Kim, W. C., and Mulvey, J. M. (2014). Dynamic asset allocation for varied financial markets under regime switching framework. *European Journal of Operational Research* **234**(2), 450–458 (<https://doi.org/10.1016/j.ejor.2013.03.032>).
- Baum, L. E., Petrie, T., Soules, G., and Weiss, N. (1970). A maximization technique occurring in the statistical analysis of probabilistic functions of Markov chains. *Annals of Mathematical Statistics* **41**(1), 164–171 (<https://doi.org/10.1214/aoms/1177697196>).

- Bellman, R. E. (1961). *Adaptive Control Processes: A Guided Tour*. Princeton University Press (<https://doi.org/10.1515/9781400874668>).
- Best, M. J., and Grauer, R. R. (1991). On the sensitivity of mean-variance-efficient portfolios to changes in asset means: some analytical and computational results. *Review of Financial Studies* **4**(2), 315–342 (<https://doi.org/10.1093/rfs/4.2.315>).
- Boyd, S. P., El Ghaoui, L., Feron, E., and Balakrishnan, V. (1994). *Linear Matrix Inequalities in System and Control Theory*. Studies in Applied and Numerical Mathematics, Volume 15. Society for Industrial and Applied Mathematics, Philadelphia, PA (<https://doi.org/10.1137/1.9781611970777>).
- Broadie, M. (1993). Computing efficient frontiers using estimated parameters. *Annals of Operations Research* **45**(1), 21–58 (<https://doi.org/10.1007/BF02282040>).
- Chopra, V. K., and Ziemba, W. T. (1993). The effect of errors in means, variances, and covariances on optimal portfolio choice. *Journal of Portfolio Management* **19**(2), 6–11 (<https://doi.org/10.3905/jpm.1993.409440>).
- Costa, G., and Kwon, R. H. (2019). Risk parity portfolio optimization under a Markov regime-switching framework. *Quantitative Finance* **19**(3), 453–471 (<https://doi.org/10.1080/14697688.2018.1486036>).
- Delage, E., and Ye, Y. (2010). Distributionally robust optimization under moment uncertainty with application to data-driven problems. *Operations Research* **58**(3), 595–612 (<https://doi.org/10.1287/opre.1090.0741>).
- Dunning, I., Huchette, J., and Lubin, M. (2017). JuMP: a modeling language for mathematical optimization. *SIAM Review* **59**(2), 295–320 (<https://doi.org/10.1137/15M1020575>).
- Fama, E. F., and French, K. R. (1993). Common risk factors in the returns on stocks and bonds. *Journal of Financial Economics* **33**(1), 3–56 ([https://doi.org/10.1016/0304-405X\(93\)90023-5](https://doi.org/10.1016/0304-405X(93)90023-5)).
- French, K. R. (2019). Current research returns. Data Library, Dartmouth College. URL: <http://mba.tuck.dartmouth.edu/pages/faculty/ken.french/data.library.html>.
- Goldfarb, D., and Iyengar, G. (2003). Robust portfolio selection problems. *Mathematics of Operations Research* **28**(1), 1–38 (<https://doi.org/10.1287/moor.28.1.1.14260>).
- Guidolin, M. (2011). Markov switching models in empirical finance. In *Missing Data Methods: Time-Series Methods and Applications*, Drukker, D. M. (ed), pp. 1–86. Advances in Econometrics, Volume 27, Part 2. Emerald Group Publishing, Bingley ([https://doi.org/10.1108/S0731-9053\(2011\)000027B004](https://doi.org/10.1108/S0731-9053(2011)000027B004)).
- Guidolin, M., and Timmermann, A. (2007). Asset allocation under multivariate regime switching. *Journal of Economic Dynamics and Control* **31**(11), 3503–3544 (<https://doi.org/10.1016/j.jedc.2006.12.004>).
- Hamilton, J. D. (1989). A new approach to the economic analysis of nonstationary time series and the business cycle. *Econometrica* **57**(2), 357–384 (<https://doi.org/10.2307/1912559>).
- Hamilton, J. D. (2010). Regime switching models. In *Macroeconometrics and Time Series Analysis*, Durlauf, S. N., and Blume, L. E. (eds), pp. 202–209. The New Palgrave Economics Collection. Palgrave Macmillan, London (https://doi.org/10.1057/9780230280830_23).
- Jorion, P. (1986). Bayes–Stein estimation for portfolio analysis. *Journal of Financial and Quantitative Analysis* **21**(3), 279–292 (<https://doi.org/10.2307/2331042>).

- Kritzman, M., Page, S., and Turkington, D. (2012). Regime shifts: implications for dynamic strategies. *Financial Analysts Journal* **68**(3), 22–39 (<https://doi.org/10.2469/faj.v68.n3.3>).
- Ledoit, O., and Wolf, M. (2003). Improved estimation of the covariance matrix of stock returns with an application to portfolio selection. *Journal of Empirical Finance* **10**(5), 603–621 ([https://doi.org/10.1016/S0927-5398\(03\)00007-0](https://doi.org/10.1016/S0927-5398(03)00007-0)).
- Levy, M., and Kaplanski, G. (2015). Portfolio selection in a two-regime world. *European Journal of Operational Research* **242**(2), 514–524 (<https://doi.org/10.1016/j.ejor.2014.10.012>).
- Lintner, J. (1965). The valuation of risk assets and the selection of risky investments in stock portfolios and capital budgets. *Review of Economics and Statistics* **47**(1), 13–37 (<https://doi.org/10.2307/1924119>).
- Liu, X., and Chen, R. (2016). Regime-switching factor models for high-dimensional time series. *Statistica Sinica* **26**(4), 1427–1451 (<https://doi.org/10.5705/ss.2014.265t>).
- Markowitz, H. (1952). Portfolio selection. *Journal of Finance* **7**(1), 77–91 (<https://doi.org/10.1111/j.1540-6261.1952.tb01525.x>).
- Merton, R. C. (1980). On estimating the expected return on the market: an exploratory investigation. *Journal of Financial Economics* **8**(4), 323–361 ([https://doi.org/10.1016/0304-405X\(80\)90007-0](https://doi.org/10.1016/0304-405X(80)90007-0)).
- Michaud, R. O., and Michaud, R. O. (2008). *Efficient Asset Management: A Practical Guide to Stock Portfolio Optimization and Asset Allocation*, 2nd edn. Oxford University Press.
- Mossin, J. (1966). Equilibrium in a capital asset market. *Econometrica* **34**(4), 768–783 (<https://doi.org/10.2307/1910098>).
- Nystrup, P., Madsen, H., and Lindström, E. (2018). Dynamic portfolio optimization across hidden market regimes. *Quantitative Finance* **18**(1), 83–95 (<https://doi.org/10.1080/14697688.2017.1342857>).
- Schwarz, G. (1978). Estimating the dimension of a model. *Annals of Statistics* **6**(2), 461–464 (<https://doi.org/10.1214/aos/1176344136>).
- Sharpe, W. F. (1964). Capital asset prices: a theory of market equilibrium under conditions of risk. *Journal of Finance* **19**(3), 425–442 (<https://doi.org/10.1111/j.1540-6261.1964.tb02865.x>).
- Sharpe, W. F. (1994). The Sharpe ratio. *Journal of Portfolio Management* **21**(1), 49–58 (<https://doi.org/10.3905/jpm.1994.409501>).
- Tütüncü, R. H., and Koenig, M. (2004). Robust asset allocation. *Annals of Operations Research* **132**, 157–187 (<https://doi.org/10.1023/B:ANOR.0000045281.41041.ed>).

